

# Toward Coordination-free and Reconfigurable Mixed Concurrency Control

Dixin Tang

Aaron J. Elmore




THE UNIVERSITY OF  
**CHICAGO**

# Hardware Development Changes Database Architecture




THE UNIVERSITY OF  
**CHICAGO**

Disk-based Database



A Few Cores




Small Memory




Disk Stalls Dominate the Performance

Large Disk




Main Memory Database



**Concurrency Control  
Becomes the New Bottleneck**



# A Closer Look at Concurrency Control



## Two Goals of Concurrency Control

- Interleaving concurrent operations to maximize the performance
- Guarantee consistency

### Serializability


The result of interleaved execution of concurrent transactions **is equivalent to** the result of executing these transactions in one serial order

# One Concurrency Control Does Not Fit All



## PartCC (Partition-based single-thread concurrency control)

- Perform well under partitionable workloads
- Cross-partition transactions hurt the performance



# One Concurrency Control Does not Fit All





## OCC (Optimistic Concurrency Control)

- ❖ Perform better under low-conflict workloads
- ❖ Conflicts can hurt the performance because transactions need to be restarted if validation fails

## 2PL (Two Phase Locking)

- ❖ Perform better under highly-conflicted workloads
- ❖ Concurrency control overhead and synchronization overhead




# One Concurrency Control Does not Fit All



experiment on our main-memory database prototype using YCSB workloads

- ❖ OCC from Silo [1]
- ❖ 2PL using VLL [2]
- ❖ PartCC from H-Store [3]




[1] TU, S., et al. Speedy transactions in multicore in-memory databases. SOSP' 13

[2] REN, K. et al. Lightweight locking for main memory database systems. PVLDB'12

[3] KALLMAN, R., et al. H-store: a high-performance, distributed main memory transaction processing system. PVLDB'08

# A General Solution: Mixed Concurrency Control




## Two Benefits

- Each protocol can process the part of workload it is optimized for
- Each protocol can avoid being brittle to workload where it does not perform well

# Two Challenges of Mixed Concurrency Control



- How to partition a workload and mix multiple concurrency control protocols efficiently
- How to reconfigure a protocol when the workload changes




# Previous Approach of Mixed Concurrency Control: Partition Stored Procedures by Conflicts



THE UNIVERSITY OF  
**CHICAGO**


**Stored Procedure (SP)**  
A Parameterized Transaction Template



## Previous Approach [1]


- ❖ Group stored procedures by conflicts extracted from their semantic information
- ❖ Assign each group a protocol to process conflicts within that group
- ❖ Need additional concurrency control protocols to process conflicts across groups

# Previous Approach of Mixed Concurrency Control: Partition Stored Procedures by Conflicts






# A New Perspective: Partition Records by Access Characteristics



# CormCC: Coordination-free and Reconfigurable Mixed Concurrency Control



**Coordination-free:** no additional concurrency control protocols

**Reconfigurable:** partitioning records depend on real-time data access characteristics, make online protocol reconfiguration possible


## Our Approach

- ❖ Partition database records and assign each partition a single protocol
- ❖ A single protocol is used to process all operations for that partition of records


To achieve the two goals, we need to answer four questions:

- ❖ How does CormCC execute
- ❖ How to maintain serializability
- ❖ How to guarantee deadlock free
- ❖ How to enable online protocol switch


# CormCC Execution



# Correctness of CormCC – Serializability



- ❖ COCSR (Commit ordering conflict serializable)
  - ❖ Sufficient condition of serializable
  - ❖ Commit ordering respects conflicts
- ❖ If all protocols are COCSR, then CormCC is COCSR
- ❖ Proof can be found in the paper




# CormCC – Deadlock Free




THE UNIVERSITY OF  
CHICAGO

## Deadlock Detection



## Our Approach: Deadlock Prevention

- ❖ We require each protocol can only exclusively let transactions wait in no more than one phase
  - ❖ No deadlock within one phase
  - ❖ Transactions in earlier phases can wait for later phases, but not the other way around




# CormCC – Online Protocol Reconfiguration



THE UNIVERSITY OF  
CHICAGO

Online reconfiguration can cause inconsistency



A straightforward solution: stop all


- ❖ Waiting all working threads to complete their current transactions, and stop them from receiving new transactions
- ❖ Decrease the performance of database

# CormCC – Online Protocol Reconfiguration



THE UNIVERSITY OF  
CHICAGO

Our solution: using a mediated protocol that is compatible to both old and new protocols



# CormCC – Online Protocol Reconfiguration

## How to Build a Mediated Protocol

- ❖ The mediated protocol executes the logics of both old and new protocol
- ❖ Example: Mediated Protocol between OCC and 2PL

## Mediated Protocol between OCC and 2PL

### Read

- ❖ Read timestamp (OCC)
- ❖ Apply read lock (2PL)

### Write

- ❖ Write to a local buffer (OCC)
- ❖ Apply write lock (2PL)


### Validate

- ❖ Execute Validate phase of OCC

# Prototype Design



- ❖ Supporting PartCC from H-Store [3], OCC from Silo [1], and 2PL from VLL [2]
- ❖ Partition the whole database and apply each partition a single protocol



[1] TU, S., et al. Speedy transactions in multicore in-memory databases. SOSP' 13


[2] REN, K. et al. Lightweight locking for main memory database systems. PVLDB'12

[3] KALLMAN, R., et al. H-store: a high-performance, distributed main memory transaction processing system. PVLDB'08

# Prototype Design



- ❖ Selecting the ideal protocol for each partition
  - ❖ Feature Engineering: design several features to capture the performance difference of candidate protocols
  - ❖ Classifiers: building a two-layer classifier



## Experiment Settings

- ❖ A Machine with 32 cores, 256 GB main memory

## Workload


- ❖ TPC-C: Order processing application with 5 stored procedures, 32 warehouses
- ❖ YCSB: One Table with 1 million records, each with 25 columns and 20 bytes for each column; One type of transaction mixed with read or read-and-modify operations
- ❖ Partition into 32 partitions (i.e. equal to the number of cores)

# Experiments

Our experiments compare CormCC with

- ❖ Single candidate protocols
- ❖ Hybrid - Hybrid OCC and 2PL execution based on MOCC [1] /Hsync [2]
- ❖ Tebaldi - A stored procedure-oriented general framework of mixed concurrency control [3]

Tebaldi Configuration for TPC-C (from Tebaldi paper)



[1] WANG, T., et al. Mostly-optimistic concurrency control for highly contended dynamic workloads on a thousand cores. PVLDB'16

[2] SHANG, Z., et al. Graph analytics through fine-grained parallelism. SIGMOD'16

[3] SU, C., et al. Bringing modular concurrency control to the next level. SIGMOD' 17


# Experiments – Compare with Tebaldi



THE UNIVERSITY OF  
CHICAGO

## Comparison under different partitionability (TPC-C)


- ❖ Start with well-partitionable workload
  - ❖ Each partition receives 100% single-partition transactions
- ❖ Increase the number of partitions receiving 100% cross-partition transactions



# Experiments – Compare with Tebaldi

## Comparison under different conflict rates (TPC-C)

- ❖ We modify TPC-C to increase access skewness
  - ❖ We use Zipfian distribution and increase its theta parameter to introduce higher conflict rate



# Experiments – Varied workloads




- ❖ Vary parameters every 5 seconds:

Transactions mix, Percentages of cross-partition transactions, Access skewness (i.e. theta of Zipf)

Tests over varied workloads (TPC-C)


—●— PartCC —●— OCC —●— 2PL —●— Hybrid —●— CormCC



## Experiments – A closer look at the same results of varied workload tests




- ❖ We aggregate the throughput of every 5s and report CormCC throughput ratio to single protocols
  - ❖ Max: Speedup over the worst single protocols
  - ❖ Min: Speedup over the best single protocols
  - ❖ Average: Speedup over average throughput of single protocols



# Experiments – Mediated Protocol Switch

- ❖ Switch from OCC to 2PL using YCSB workloads
- ❖ Report the throughput during protocol switch
- ❖ Compared with StopALL
- ❖ Test a workload with short-only transactions and workloads with one long transaction of different duration
  - ❖ 0.5s, 1s, 2s, 4s
- ❖ Test different switching points
  - ❖ At the **beginning** of the long transaction
  - ❖ At the **middle** of the long transaction
  - ❖ At the **end** of the long transaction



# Conclusion



THE UNIVERSITY OF  
CHICAGO


- ❖ CormCC, a general mixed concurrency control framework that does not introduce any coordination overhead and supports online reconfiguration
- ❖ Experiments show that CormCC can achieve significant throughput improvement over single static protocols and state-of-the-art mixed approaches.

Thanks!



# Backup Slides

# CormCC Execution in Prototype




# Experiments – Varied workloads (YCSB)



THE UNIVERSITY OF  
CHICAGO

- ❖ Vary parameters every 5 seconds:

Transactions mix, Percentages of cross-partition transactions, Access skewness (i.e. theta of Zipf)



# Experiments – Aggregated Results of Varied workloads (YCSB)



THE UNIVERSITY OF  
CHICAGO

